



Linear Dimensionality Reduction: PCA

Qinliang Su (苏勤亮)

Sun Yat-sen University

suqliang@mail.sysu.edu.cn

Outline

- Motivation
- Perspective 1: Minimizing Reconstruction Error
- Perspective 2: Maximizing Variance
- Perspective 3: SVD
- Other Applications of PCA

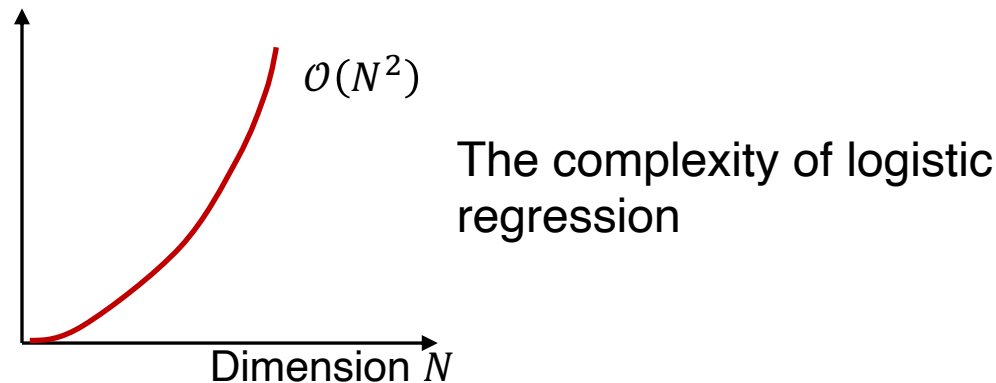
Motivation

- The dimensionality of many types of data is very high, *e.g.*, the dimension of images below is as high as

$$256 \times 256 = 65536$$




- If we work on the raw data directly, the complexity of subsequent tasks (*e.g.*, classification) could be extremely high




Why the dimensionality could be reduced?

- The high-dimensional data often resides on a low-dimensional intrinsic space approximately



3-dimensional data lies on a 2-dimensional plane approximately




2-dimensional data lies on a 1-dimensional line approximately

The key is how to find *the principal directions* under which the dimensions of data can be significantly reduced

- For the real-world data, this is also possible to find the low-dimensional space

e.g., images of human face can be well represented *with only several values* if appropriate directions can be found



$$x \approx \mu_0 + a_1\mu_1 + \cdots + a_7\mu_7$$

The raw image x that has 65536 values can be represented by only 7 values of $a_1, \cdots a_7$

Outline

- Motivation
- Perspective 1: Minimizing the Reconstruction Error
- Perspective 2: Maximizing Variance
- Perspective 3: SVD
- Other Applications of PCA

Re-representation under New Directions

- How to represent **orthogonal directions** in high dimensional space?

A set of vectors \mathbf{u}_i satisfying

$$\mathbf{u}_i^T \mathbf{u}_j = \delta_{ij}$$

where $\delta_{ij} = 1$ if $i = j$; 0 otherwise

Theorem: Under the M given orthogonal directions $\{\mathbf{u}_i\}_{i=1}^M$, the **best approximation** to a data sample \mathbf{x} is

$$\tilde{\mathbf{x}} = \alpha_1 \mathbf{u}_1 + \alpha_2 \mathbf{u}_2 + \cdots + \alpha_M \mathbf{u}_M$$

with the α_i equal to

$$\alpha_i = \mathbf{u}_i^T \mathbf{x}$$

Proof:

$$\begin{aligned}\|\mathbf{x} - \tilde{\mathbf{x}}\|^2 &= \left\| \mathbf{x} - \sum_{i=1}^M \alpha_i \mathbf{u}_i \right\|^2 \\ &= \|\mathbf{x}\|^2 - 2 \sum_{i=1}^M \alpha_i \mathbf{u}_i^T \mathbf{x} + \sum_{i=1}^M \alpha_i^2\end{aligned}$$

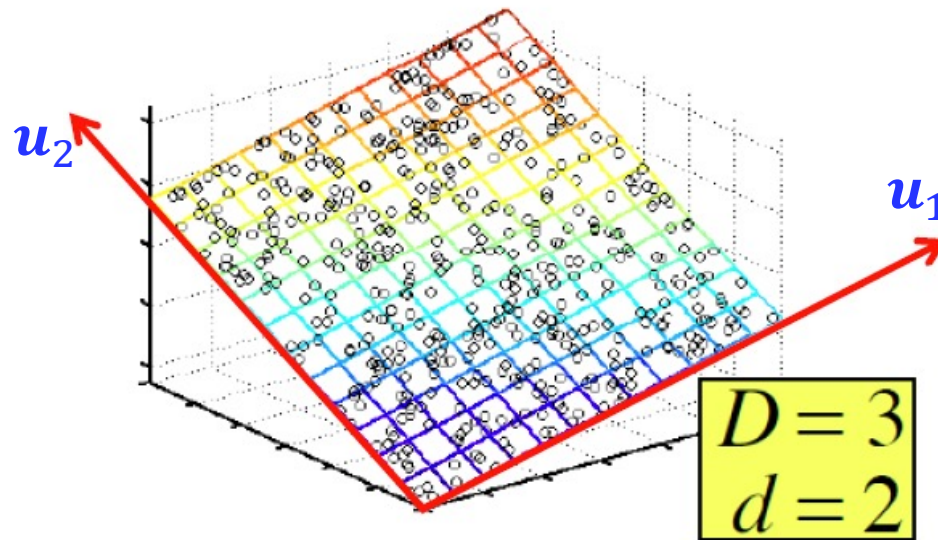
where we used $\mathbf{u}_i^T \mathbf{u}_j = 0$ for $i \neq j$ and 1 for $i = j$

This is a quadratic function, and can be minimized when $\alpha_i = \mathbf{u}_i^T \mathbf{x}$

Given the directions $\{\mathbf{u}_i\}_{i=1}^M$, the best coefficient is $\alpha_i = \mathbf{u}_i^T \mathbf{x}$.
But *how to find the best directions is still unknown*

Finding the Best Directions

- **Goal:** Given data samples $\{\mathbf{x}^{(n)}\}_{n=1}^N$ from \mathbb{R}^D , finding M orthogonal directions \mathbf{u}_i so that the original data can be best represented under them



$$\mathbf{x}^{(n)} \approx \sum_{i=1}^M \alpha_i^{(n)} \mathbf{u}_i$$

- Suppose the best directions $\{\mathbf{u}_i\}_{i=1}^M$ are given, what are the best coefficients $\alpha_i^{(n)}$?

$$\alpha_i^{(n)} = \mathbf{u}_i^T \mathbf{x}^{(n)}$$

Instead of representing the data $\mathbf{x}^{(n)}$ directly, we first center the data to the origin, *i.e.*, subtracting each data point $\mathbf{x}^{(n)}$ by its mean

$$\mathbf{x}^{(n)} - \bar{\mathbf{x}}, \quad \text{where } \bar{\mathbf{x}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)}$$

- The objective can now be described as minimizing the error between $\mathbf{x}^{(n)} - \bar{\mathbf{x}}$ and its best approximant $\tilde{\mathbf{x}}^{(n)} = \sum_{i=1}^M \alpha_i^{(n)} \mathbf{u}_i$ in the $\text{span}(\mathbf{u}_1, \dots, \mathbf{u}_M)$

$$E = \frac{1}{N} \sum_{n=1}^N \left\| (\mathbf{x}^{(n)} - \bar{\mathbf{x}}) - \tilde{\mathbf{x}}^{(n)} \right\|^2$$

where the best coefficient α_i is known equal to

$$\alpha_i = \mathbf{u}_i^T (\mathbf{x}^{(n)} - \bar{\mathbf{x}})$$

- Reformulating the reconstruction error E

a) Substituting $\tilde{\mathbf{x}}^{(n)} = \sum_{i=1}^M \alpha_i^{(n)} \mathbf{u}_i$ into $E = \frac{1}{N} \sum_{n=1}^N \|(\mathbf{x}^{(n)} - \bar{\mathbf{x}}) - \tilde{\mathbf{x}}^{(n)}\|^2$ and using $\mathbf{u}_i^T \mathbf{u}_j = \delta_{ij}$ gives

$$E = \frac{1}{N} \left(\sum_{n=1}^N \|\mathbf{x}^{(n)} - \bar{\mathbf{x}}\|^2 - 2 \sum_{n=1}^N \sum_{i=1}^M \alpha_i^{(n)} (\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T \mathbf{u}_i + \sum_{n=1}^N \sum_{i=1}^M (\alpha_i^{(n)})^2 \right)$$

b) Substituting $\alpha_i^{(n)} = \mathbf{u}_i^T (\mathbf{x}^{(n)} - \bar{\mathbf{x}})$ gives

$$E = \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}^{(n)} - \bar{\mathbf{x}}\|^2 - \sum_{i=1}^M \mathbf{u}_i^T \underbrace{\frac{1}{N} \sum_{n=1}^N (\mathbf{x}^{(n)} - \bar{\mathbf{x}})(\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T}_{\triangleq \mathbf{S}} \mathbf{u}_i$$

c) Writing it into a matrix form gives

$$E = \frac{1}{N} \|\mathbf{X} - \bar{\mathbf{X}}\|_F^2 - \frac{1}{N} \sum_{i=1}^M \mathbf{u}_i^T \mathbf{S} \mathbf{u}_i$$

where $\mathbf{X} \triangleq [\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}]$ and $\|\cdot\|_F$ is the Frobenius norm

- Minimizing $E = \frac{1}{N} \|\mathbf{X} - \bar{\mathbf{X}}\|_F^2 - \frac{1}{N} \sum_{i=1}^M \mathbf{u}_i^T \mathbf{S} \mathbf{u}_i$ under the constraint $\mathbf{u}_i^T \mathbf{u}_j = \delta_{ij}$ is equivalent to maximize

$$\max_{\mathbf{u}_1 \cdots \mathbf{u}_M} \sum_{i=1}^M \mathbf{u}_i^T \mathbf{S} \mathbf{u}_i$$

$$s. t. : \mathbf{u}_i^T \mathbf{u}_j = \delta_{ij}$$

- Consider the simple case with $M = 1$. The problem is reduced to:

$$\begin{aligned} \max_{\mathbf{u}_1} \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 \\ \text{s.t.} : \mathbf{u}_1^T \mathbf{u}_1 = 1 \end{aligned}$$

- This is equivalent to maximize

$$\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 - \lambda(\mathbf{u}_1^T \mathbf{u}_1 - 1)$$

- Taking the derivative *w.r.t.* \mathbf{u}_1 and setting it to 0 gives

$$\mathbf{S} \mathbf{u}_1 = \lambda \mathbf{u}_1,$$

from which we can see that \mathbf{u}_1 should be **the eigenvector of \mathbf{S}**

- It can be further checked that it is **the eigenvector *w.r.t. to the largest eigenvalue* that maximizes $\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1$**

- For the case of $M = 2$, the problem becomes

$$\max_{\mathbf{u}_1, \mathbf{u}_2} \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 + \mathbf{u}_2^T \mathbf{S} \mathbf{u}_2$$

$$s.t.: \mathbf{u}_1^T \mathbf{u}_1 = 1, \mathbf{u}_2^T \mathbf{u}_2 = 1, \mathbf{u}_1^T \mathbf{u}_2 = 0$$

- This is equivalent to maximize

$$\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 - \lambda_1 (\mathbf{u}_1^T \mathbf{u}_1 - 1) + \mathbf{u}_2^T \mathbf{S} \mathbf{u}_2 - \lambda_2 (\mathbf{u}_2^T \mathbf{u}_2 - 1)$$

under the constraint $\mathbf{u}_1^T \mathbf{u}_2 = 0$

- Taking the derivative w.r.t. \mathbf{u}_1 and \mathbf{u}_2 and setting it to 0 gives

$$\mathbf{S} \mathbf{u}_1 = \lambda_1 \mathbf{u}_1, \quad \mathbf{S} \mathbf{u}_2 = \lambda_2 \mathbf{u}_2,$$

⇒ \mathbf{u}_1 and \mathbf{u}_2 must be the **eigenvectors of \mathbf{S}**

⇒ It can be seen that to maximize $\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 + \mathbf{u}_2^T \mathbf{S} \mathbf{u}_2$, \mathbf{u}_1 and \mathbf{u}_2 must be the eigenvectors **corresponding to the two largest eigenvalues**

For the case $M > 1$, the directions \mathbf{u}_i are *the eigenvectors of S corresponding to the largest M eigenvalues*

$$\mathbf{A} = \mathbf{X} - \bar{\mathbf{X}}$$

Question: Will the eigenvectors \mathbf{u}_i of S satisfy $\mathbf{u}_i^T \mathbf{u}_j = 0$ for $i \neq j$?

- For any $D \times D$ *real symmetric* matrix like $S \triangleq \mathbf{A}\mathbf{A}^T$, it has D eigenvectors that are orthogonal to each other
- For every $S \triangleq \mathbf{A}\mathbf{A}^T$, it can be decomposed as

$$\mathbf{S} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T$$

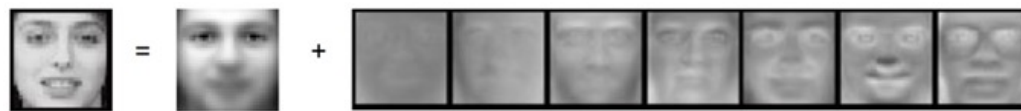
where \mathbf{U} is comprised of the eigenvectors of S and $\mathbf{U}\mathbf{U}^T = \mathbf{I}$;
 $\mathbf{\Lambda}$ is a diagonal matrix composed of eigenvalues of S

Examples

Input data: each face image is a data point



Top 25 principal directions



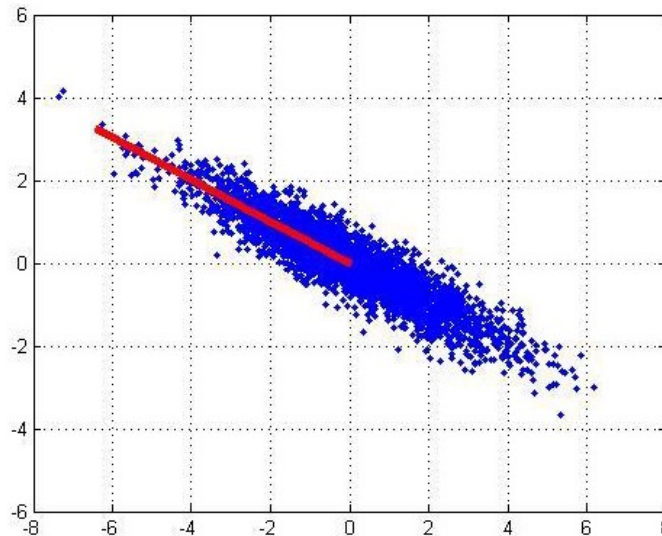
$$x \approx \bar{x} + \alpha_1 u_1 + \cdots + \alpha_7 u_7$$

Outline

- Motivation
- Perspective 1: Minimizing the Reconstruction Error
- **Perspective 2: Maximizing Variance**
- Perspective 3: SVD
- Other Applications of PCA

Problem Formulation

- **Goal:** Given a dataset $\{\mathbf{x}^{(n)}\}_{n=1}^N$, find orthogonal directions $\{\mathbf{u}_k\}_{k=1}^M$ such that **the variance** of data projected onto these directions are maximized



Maximizing the variance is equivalent to *preserve the information of original data as much as possible*

- For the first direction \mathbf{u}_1 , we hope the variance of projected data along the direction \mathbf{u}_1 , i.e., $\{\mathbf{u}_1^T \mathbf{x}^{(n)}\}_{n=1}^N$, is maximized

➤ The variance expression

$$\begin{aligned} var &= \frac{1}{N} \sum_{n=1}^N \left(\mathbf{u}_1^T (\mathbf{x}^{(n)} - \bar{\mathbf{x}}) \right)^2 \\ &= \mathbf{u}_1^T \frac{1}{N} \sum_{n=1}^N (\mathbf{x}^{(n)} - \bar{\mathbf{x}})(\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T \mathbf{u}_1 \\ &= \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 \end{aligned}$$

- Subjecting to $\mathbf{u}_1^T \mathbf{u}_1 = 1$, as proved previously, the variance is maximized when \mathbf{u}_1 is *the eigenvector of \mathbf{S} corresponding to the largest eigenvalue*

- For the second direction \mathbf{u}_2 , it also should maximize the variance

$$var = \mathbf{u}_2^T \mathbf{S} \mathbf{u}_2,$$

but should subject to the constraints $\mathbf{u}_i^T \mathbf{u}_j = \delta_{ij}$, that is,

$$\mathbf{u}_2^T \mathbf{u}_2 = 1 \quad \mathbf{u}_1^T \mathbf{u}_2 = 0$$

- Due to \mathbf{u}_1 being the eigenvector w.r.t. the largest eigenvalue, it can be proved that *\mathbf{u}_2 is the eigenvector of \mathbf{S} corresponding the second largest eigenvalue*

\mathbf{u}_i is the eigenvector of $\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}^{(n)} - \bar{\mathbf{x}})(\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T$ corresponding the i -th largest eigenvalue

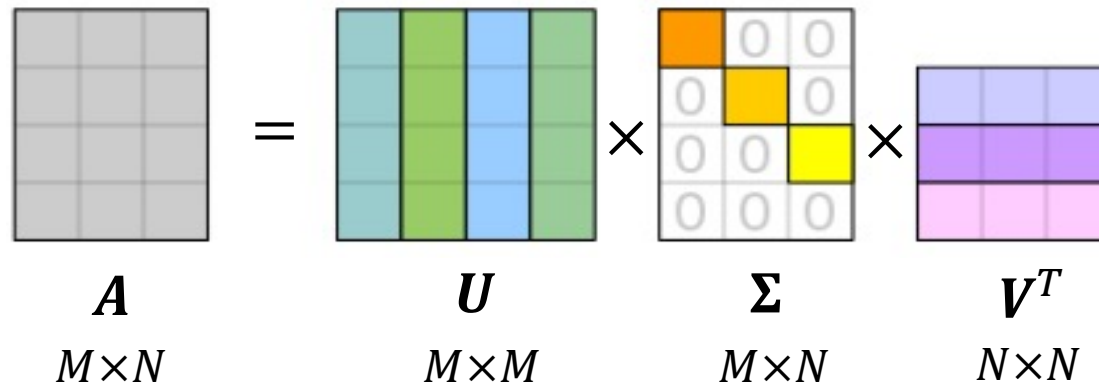
Outline

- Motivation
- Perspective 1: Minimizing Reconstruction Error
- Perspective 2: Maximizing Variance
- **Perspective 3: SVD**
- Other Applications of PCA

Singular Value Decomposition (SVD)

- For any $M \times N$ matrix A , it can always be decomposed as

$$A = U \Sigma V^T$$



- $U = [\mathbf{u}_1, \dots, \mathbf{u}_M]$ and $V = [\mathbf{v}_1, \dots, \mathbf{v}_N]$, with \mathbf{u}_i and \mathbf{v}_i being the i -th eigenvector of AA^T and $A^T A$, and $\mathbf{u}_i^T \mathbf{u}_j = \delta_{ij}$ and $\mathbf{v}_i^T \mathbf{v}_j = \delta_{ij}$
- Σ only has nonzero values on the diagonal, which are the squared root of eigenvalues of AA^T or $A^T A$ (Nonzero eigenvalues of AA^T and $A^T A$ are the same)

Σ_{ii} is called *singular values* and are stored in a descending order

- Because Σ only has nonzero values on the diagonal, A can be expressed as

$$A = U' \Sigma' V'^T = \sum_{i=1}^r \Sigma_{ii} \mathbf{u}_i \mathbf{v}_i^T$$

where \mathbf{u}_i and \mathbf{v}_i are the i -th column of U and V ; r is the number of nonzero diagonal elements in Σ

$$\begin{array}{ccccccc}
 \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline \end{array} & = & \begin{array}{|c|c|} \hline & \\ \hline & \\ \hline & \\ \hline & \\ \hline & \\ \hline \end{array} & \times & \begin{array}{|c|c|c|} \hline & 0 & \\ \hline 0 & & \\ \hline & & \\ \hline \end{array} & \times & \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline & & & & \\ \hline \end{array} \\
 A & & U' & & \Sigma' & & V'^T \\
 M \times N & & M \times r & & r \times r & & r \times N
 \end{array}$$

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 3 & 3 & 3 & 0 & 0 \\ 4 & 4 & 4 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 2 & 0 & 4 & 4 \\ 0 & 0 & 0 & 5 & 5 \\ 0 & 1 & 0 & 2 & 2 \end{bmatrix} = \begin{bmatrix} 0.13 & 0.02 & -0.01 \\ 0.41 & 0.07 & -0.03 \\ 0.55 & 0.09 & -0.04 \\ 0.68 & 0.11 & -0.05 \\ 0.15 & -0.59 & 0.65 \\ 0.07 & -0.73 & -0.67 \\ 0.07 & -0.29 & 0.32 \end{bmatrix} \begin{bmatrix} 12.4 & 0 & 0 \\ 0 & 9.5 & 0 \\ 0 & 0 & 1.3 \end{bmatrix} \begin{bmatrix} 0.56 & 0.59 & 0.56 & 0.09 & 0.09 \\ 0.12 & -0.02 & 0.12 & -0.69 & -0.69 \\ 0.40 & -0.80 & 0.40 & 0.09 & 0.09 \end{bmatrix}$$

- The vector \mathbf{u}_i in the SVD decomposition of \mathbf{A} is the eigenvector of $\mathbf{A}\mathbf{A}^T$ w.r.t. its i -th largest eigenvalues
- By setting $\mathbf{A} = \tilde{\mathbf{X}}$ with $\tilde{\mathbf{X}} \triangleq [\mathbf{x}^{(1)} - \bar{\mathbf{x}}, \mathbf{x}^{(2)} - \bar{\mathbf{x}}, \dots, \mathbf{x}^{(N)} - \bar{\mathbf{x}}]$, we can see that

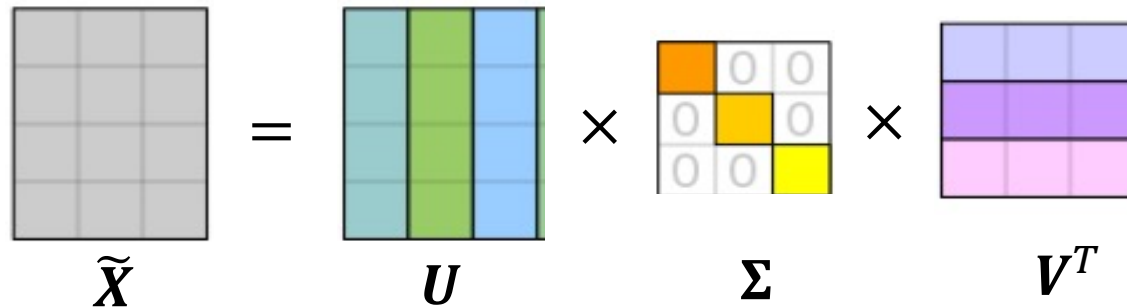
$$\begin{aligned}\mathbf{A}\mathbf{A}^T &= \sum_{n=1}^N (\mathbf{x}^{(n)} - \bar{\mathbf{x}})(\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T \\ &= \mathbf{N} \cdot \mathbf{S}\end{aligned}$$

⇒ The eigenvectors of $\mathbf{A}\mathbf{A}^T$ are the same as the matrix \mathbf{S}

By performing SVD on $\tilde{\mathbf{X}}$, principal directions of data $\{\mathbf{x}^{(n)}\}_{n=1}^N$ can be directly obtained, which are the columns of left SVD matrix \mathbf{U}

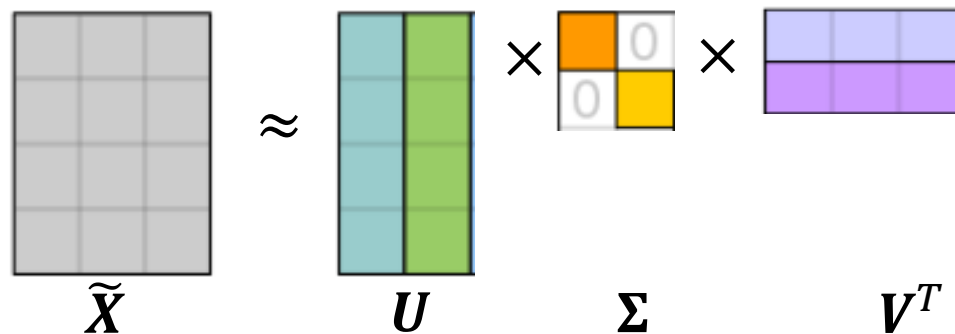
- Question:

Given the SVD decomposition of \tilde{X} as shown below, what are the principle directions and the coefficients α_i 's for $\tilde{x}^{(n)}$?



The diagram shows the SVD decomposition of a 4x4 matrix \tilde{X} (gray grid) into three matrices: U (4x3, with columns colored teal, green, and blue), Σ (3x3, diagonal with orange, yellow, and light yellow squares), and V^T (3x4, with rows colored light blue, purple, and pink). The equation is $\tilde{X} = U \Sigma V^T$.

If only top two directions are kept, what are the coefficients α_i 's?



The diagram shows the approximate SVD decomposition of \tilde{X} using only the top two principal components. \tilde{X} (4x4 gray grid) is approximated by U (4x2, with columns colored teal and green), Σ (2x2, diagonal with orange and yellow squares), and V^T (2x4, with rows colored light blue and purple). The equation is $\tilde{X} \approx U \Sigma V^T$.

Outline

- Motivation
- Perspective 1: Minimizing Reconstruction Error
- Perspective 2: Maximizing Variance
- Perspective 3: SVD
- Other Applications of PCA

Image Compression

Partition a 372×492 image below into many 12×12 patches

- Each patch is viewed as a data instance
- Performing PCA on the patches



Reconstruction Error vs # PCA components

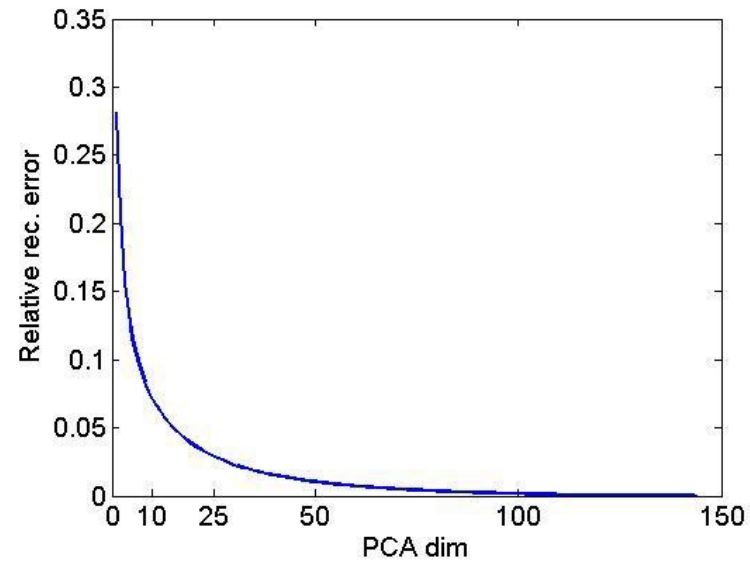
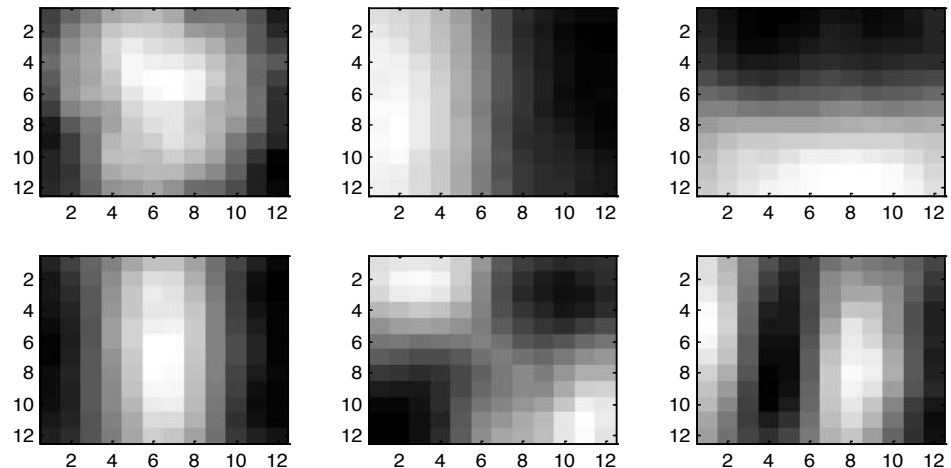


Illustration of the top 6 PCA components





Reconstruction with the top 60 components



Reconstruction with the top 16 components

Denoising

Noisy Image



Denoised Image



Reconstructed from the top 15 principal components